

Randomized Pursuit–Evasion in a Polygonal Environment

Volkan Isler, *Student Member, IEEE*, Sampath Kannan, and Sanjeev Khanna

Abstract—This paper contains two main results. First, we revisit the well-known *visibility-based pursuit–evasion* problem, and show that in contrast to deterministic strategies, a single pursuer can locate an unpredictable evader in any simply connected polygonal environment, using a randomized strategy. The evader can be arbitrarily faster than the pursuer, and it may know the position of the pursuer at all times, but it does not have prior knowledge of the random decisions made by the pursuer. Second, using the randomized algorithm, together with the solution to a problem called the “lion and man problem” as subroutines, we present a strategy for two pursuers (one of which is at least as fast as the evader) to quickly capture an evader in a simply connected polygonal environment. We show how this strategy can be extended to obtain a strategy for a polygonal room with a door, two pursuers who have only line-of-sight communication, and a single pursuer (at the expense of increased capture time).

Index Terms—Dynamic noncooperative game theory, path planning, pursuit–evasion games, randomized algorithms.

I. INTRODUCTION

PURSUIT–EVASION games are among the fundamental problems studied by robotics researchers. In a pursuit–evasion game, one or more pursuers try to capture an evader who, in turn, tries to avoid capture indefinitely. A typical example is the homicidal chauffeur game, where a driver wants to collide with a pedestrian, and the goal is to determine conditions under which he can (not) do so. Among the numerous applications of this game are collision avoidance and air traffic control.

Recently, there has been increasing interest in developing pursuit strategies (which incorporate sensing limitations) to capture intelligent evaders contaminating a complex environment [3]–[5]. The main ingredient of a pursuit–evasion game is the presence of an adversarial evader who actively avoids capture. Due to this aspect, a pursuit strategy is usually different from

a search strategy, where the target’s motion is independent of the pursuer’s (e.g., [6], [7]). Obtaining such pursuit strategies is important for surveillance applications, where we would like to locate, and perhaps capture, intruders who may be adversarial. Another application is a search-and-rescue operation, where we would like to save a victim. In this setting, even though the victim is not adversarial, a pursuit strategy is still desirable, as it guarantees a rescue, regardless of the victim’s actions.

To model the adversarial nature of the game, pursuit–evasion games are usually studied in a game-theoretic framework [8], [9]. The conditions under which the pursuer can capture the evader are obtained by studying a Hamilton–Jacobi–Isaacs equation which brings together the system equations of the pursuer and the evader. This approach has the advantage of yielding a closed-form solution of the game. Unfortunately, as the environments get complicated, solving Hamilton–Jacobi–Isaacs equations becomes intractable. Therefore, solutions of pursuit–evasion games in complex environments are usually algorithmic.

Perhaps the most well-understood game in this context is the *visibility-based pursuit–evasion game*, where one or more pursuers try to locate an evader in a polygonal environment [10]–[12]. In this game, the evader is very powerful: it has unbounded speed and global visibility, meaning that it knows the location of the pursuers at all times. In [4], the authors study a similar game in a probabilistic framework (where the evader performs a random walk), and propose a greedy algorithm.

In this paper, we propose randomized pursuer strategies for the visibility-based pursuit–evasion problem. Randomization is a powerful technique which allows us to solve many problems that are not solvable by deterministic algorithms, and has found widespread applications in many areas, ranging from computational geometry to cryptography.

As we show in the following sections, it turns out that randomization provides a drastic increase in the power of the pursuers. For example, it is known that there are simply connected environments where $\Theta(\log n)$ pursuers are required [12] in order to locate the evader with deterministic strategies. Here, n denotes the number of vertices of the polygon. In contrast, we show that a single pursuer can locate the evader in any simply connected environment with high probability, even if the evader knows the pursuer’s location at all times and has unbounded speed (*Theorem 2*). The power of randomized strategies comes from the fact that the evader has no prior knowledge of the random decisions inherent in such strategies. It is worth noting that randomized strategies work against any evader strategy, and require no prior information about the strategy of the evader.

Manuscript received July 27, 2004; revised December 22, 2004. This paper was recommended for publication by Associate Editor D. Fox and Editor I. Walker upon evaluation of the reviewers’ comments. The work of V. Isler was supported in part by MURI under DAAH-19-02-1-03-83. The work of S. Kannan was supported in part by the National Science Foundation under Grant CCR0105337 and in part by the Army Research Office under Grant DAAD19-01-1-0473. The work of S. Khanna was supported in part by an Alfred P. Sloan Research Fellowship and in part by the NSF under Career Award CCR-0093117. This paper was presented in part at the Workshop on Algorithmic Foundations of Robotics, Utrecht/Zeist, The Netherlands, July 2004.

V. Isler is with the Center for Information Technology Research in the Interest of Society (CITRIS), University of California at Berkeley, Berkeley, CA 94720 USA (e-mail: isler@eecs.berkeley.edu).

S. Kannan and S. Khanna are with the Department of Computer and Information Science, University of Pennsylvania, Philadelphia, PA 19104 USA (e-mail: kannan@cis.upenn.edu; sanjeev@cis.upenn.edu).

Digital Object Identifier 10.1109/TRO.2005.851373




Fig. 1. Single pursuer cannot capture an evader using deterministic strategies.

We also address the harder task of capturing the evader. For this problem, we present a strategy for two pursuers, one of which is at least as fast as the evader. The strategy is based on the randomized strategy to locate the evader and the known solution to a problem called the “lion and man problem” [2], which is reviewed in Section III-A. The same strategy can be used to capture the evader while protecting a door. This problem was introduced in [13] to model scenarios where the goal is to locate the evader, which may leave the polygonal area through a door and win the game.

The two-pursuer strategy can be modified so that a single pursuer can also capture the evader. However, the expected time-to-capture in this case, though finite, may be significantly longer than the expected time-to-capture with two pursuers.

Organization of the Paper: We start the paper with a motivating example for randomized strategies (Section I-A). We present preliminary concepts and definitions in Section I-B. In Section II, we address the problem of locating a fast, unpredictable evader with global visibility.

Next, in Section III, we address the task of capturing the evader in a simply connected environment. For this problem, we present a randomized strategy for two pursuers, who can communicate at all times, to quickly capture the evader. We show how this strategy can be modified for a single pursuer, at the expense of increasing the capture time, in Section IV-A. We also present extensions of the basic two-pursuer strategy for the case where the pursuers have limited communication (Section IV-B), and for a scenario where the polygonal room has a door through which the evader can escape (Section IV-C).

A. Randomized Strategies

The power of randomization in the context of pursuit–evasion games is nicely illustrated by the example in Fig. 1. A similar example can be found in [14].

In this example, a single pursuer \mathcal{P} can never locate the evader using a deterministic strategy. Let us distinguish four regions A, B, C , and D , as shown in the figure. Now suppose the pursuer has a deterministic strategy of visiting these regions in the order A, B, C, D . In this case, the evader \mathcal{E} can first hide at B and escape to D while the pursuer \mathcal{P} is visiting A . Afterwards, it can repeat the same strategy and escape to B while \mathcal{P} is at C . If \mathcal{P} visits the regions in a different order, it is easy to see that \mathcal{E} can find a similar strategy to avoid \mathcal{P} . Therefore, in this polygon, one pursuer can never locate the evader.

An alternative interpretation of this situation is the following. Suppose the polygon in Fig. 1 is contaminated with many

evaders executing all possible evader strategies. There is no deterministic pursuer strategy that guarantees that all the evaders will be caught; for any given deterministic pursuer strategy, there will be at least one evader which can avoid being located forever.

Now consider the following randomized strategy. Instead of committing to a deterministic strategy, \mathcal{P} moves to the center of the polygon and selects one of the regions $\{A, B, C, D\}$ uniformly at random and visits it. It is easy to see that if \mathcal{P} guesses the region where \mathcal{E} is located correctly, then \mathcal{E} cannot escape, and the probability of this desired event is $(1/4)$. The crucial observation is that since \mathcal{E} does not know which region \mathcal{P} will visit, it cannot choose a strategy based on the order of points visited by \mathcal{P} .

The probability of locating the evader can be made arbitrarily small by repeating the same strategy a few times. If k is the number of trials, the probability of missing in all k trials is $(3/4)^k$ in this example, which decreases exponentially with k . In general, if the probability of capture is p , the expected number of rounds to capture is $(1/p)$. Note that each round is independent. We can obtain the expected time to locate the evader as follows. Since the length of a round is bounded by the time to travel between the two furthest points in the polygon (say, T), the expected time to capture is (T/p) . By repeating the experiment roughly $(1/p) \log(1/p)$ times, we can show (using the Chernoff bound) that the pursuer has a high probability of locating the evader. For details of this analysis, the reader is referred to [15].

B. Preliminaries

Let P be the input polygon, including its interior, and V be the set of vertices of P . The letter n denotes the number of vertices of the polygon. Two points $u, v \in P$ can see each other if the line segment uv lies entirely in P .

We use $d(u, v)$ to denote the length of the shortest path from u to v that remains inside P . The shortest path has the following property.

Property 1: The shortest path between any two points u and v inside a polygon P is a polygonal path whose inner vertices are vertices of P .

The *shortest path tree from a point x in P* is defined as $\cup_{v \in V} \pi(x, v)$, where $\pi(x, v)$ denotes the shortest path from x to v . A polygon is *simply connected* if any simple closed curve inside the polygon can be shrunk to a point. In other words, a simply connected polygon does not contain any “holes.” All the polygons considered in this paper are simply connected.

The *triangulation* of a polygon is a decomposition of the polygon into triangles by a maximal set of nonintersecting diagonals (see Fig. 2). The dual of a triangulation is a graph whose vertices correspond to the triangles. There is an edge between two vertices if the corresponding triangles share a side. It is well known that the triangulation of a simply connected polygon has exactly $n - 2$ triangles. In addition, the dual of the triangulation is a tree [16].

Game Formulations: In this paper, we study two pursuit–evasion games with different objectives. Both games take place in a simply connected polygon P , which is known to all players.




Fig. 2. Triangulation of a polygon and its dual tree.

The first game, which we call the *locating game*,¹ is defined as follows.

It is played between an evader and a single pursuer. An *evader trajectory* is a continuous function $e : [0, \infty) \rightarrow P$, such that $e(t)$ denotes the evader's position at time t . The *pursuer trajectory*, $p(t)$, is defined similarly. The pursuer moves with unit speed so that $\| \dot{p} \| = 1$. The *diameter* of the polygon P , denoted $\text{diam}(P)$, is defined as $\max_{u,v \in P} d(u,v)$. Since the pursuer moves with unit speed, the diameter is also equal to the maximum amount of time it takes the pursuer to travel between two points in P . The evader can be arbitrarily faster than the pursuer, but it must move continuously.

After the game starts, the players can observe their surroundings continuously. Further, at any given time t , the pursuer's location $p(t)$ is revealed to the evader. Therefore, the evader's strategy is a function of both the environment and the pursuer's trajectory. Since the pursuer does not observe the evader until the end of the game, a pursuer strategy is a function of only the environment. The pursuer wins the game if, in finite time t^* , he can reach a position such that $p(t^*)$ sees $e(t^*)$. The evader wins the game otherwise, i.e., if for any given pursuer strategy, there exists a strategy for the evader to avoid being seen by a pursuer forever.

We assume that in both games, the evader knows the strategy of the pursuer(s) before the game starts. However, it does not have access to the outcome of the random coin tosses during the execution of the pursuer's strategy. The pursuer, on the other hand, knows nothing about the evader's strategy.

The second game is called the *capture game* and is defined as follows. Let $e(t)$ denote the evader's and $p_i(t)$ denote the i th pursuer's trajectories, as before. Instead of finding the evader, the pursuers win the capture game if in finite time t^* , they can reach a position such that there exists an i with $p_i(t^*) = e(t^*)$.

In this game, one of the pursuers is as fast as the evader. Without loss of generality, we assume that $\| \dot{e} \| = \| \dot{p}_1 \| = 1$. Similar to the locating game, the players observe their surroundings continuously, and at any given time t , the pursuers' location $p_i(t)$ is revealed to the evader.

Unlike the previous game, we assume that the players move in discrete time intervals and in turns; the evader first, followed




Fig. 3. Intuitive explanation of the pursuer strategy. Each component is displayed with a different color. While the pursuer is at Δ , the evader cannot move from one component to another. Further, when the pursuer moves to Δ' , he can restrict the game to a smaller polygon.

by the pursuers. It is easy to see that a pursuit strategy that captures the evader in this formulation can be modified to a pursuit strategy which guarantees that a pursuer p_i can reach a point within unit distance from the evader in a game where the players move continuously and simultaneously. It is interesting to note that for 25 years, the basic lion's strategy for the discrete-time formulation (Section III-A) was believed to be sufficient to capture the evader in the continuous formulation, as well. However, in 1952, Besicovitch showed that this is incorrect and that the evader can escape. In [17], Littlewood shows how this result can be generalized, and that it is not possible to capture the evader in the continuous formulation (see also [18]).

II. LOCATING THE EVADER

In this section, we study the locating game and show that for any simple polygon P , the pursuer can locate the evader in $O(n \cdot \text{diam}(P))$ expected time.

A. The Pursuer Strategy

The pursuer strategy to locate the evader uses the acyclic structure of the triangulation dual of a simply connected polygon. Intuitively, it relies on the the following observation. Let Δ be the triangle that contains the pursuer's current location (see Fig. 3), and suppose Δ is nonleaf, i.e., has more than one neighbor. Since the triangulation dual is a tree, it is easy to see that Δ is a separator. Removing it from the polygon results in smaller, disconnected polygons, called components. This implies that the evader cannot move from one component to another while the pursuer is located at Δ without revealing itself.

The second observation is that the pursuer can not only prevent the evader from moving between components, but also restrict the game to a smaller polygon. While the pursuer is at Δ , let C_1, C_2 , and C_3 be the components, such that the evader is located in C_1 . Let Δ' be the neighbor of Δ in C_1 . If the pursuer moves to Δ' , he can restrict the game to C_1 , as the evader will not be able to move to any triangle not contained in C_1 .

¹The general version of this game is known as the visibility-based pursuit-evasion game [12].




Fig. 4. Choosing the next triangle within a round (Lemma 1).

Therefore, had the pursuer known the subtree that contains the evader, he could gradually move toward it, trapping the evader in smaller and smaller components. This process guarantees that the pursuer can enter the triangle which contains the evader, and this clearly implies that the evader would be located.

The main difficulty in implementing the strategy above is that the pursuer does not know the component that contains the evader. In what follows, we show how the pursuer can use randomization to tackle this difficulty.

The pursuer strategy is divided into rounds. Each round lasts at most $\text{diam}(P)$ time units.² The pursuer will start the round at a leaf triangle and end the round at another leaf triangle.³

Let Δ_0 be the pursuer's triangle at the beginning of the round and T be the triangulation tree (see Fig. 2) rooted at Δ_0 . Within the round, the pursuer will visit triangles $\Delta_0, \Delta_1, \dots$, where Δ_{i+1} is chosen from among the children of Δ_i as follows. When the pursuer is at Δ_i , let $\Delta_i^1, \dots, \Delta_i^k$ be the immediate children of Δ_i (see Fig. 4). For each child Δ_i^j , let l_i^j denote the number of leaves of the subtree rooted at Δ_i^j . Let $L = \sum_{j=1}^k l_i^j$. The next triangle, Δ_{i+1} , is chosen randomly from among the children Δ_i^j according to the following distribution. The probability that Δ_i^j is chosen for Δ_{i+1} is (l_i^j/L) . After choosing the next triangle Δ_{i+1} , the pursuer moves there. If he arrives at a leaf triangle, the round is over. Otherwise, the pursuer continues the round by picking one of the children of Δ_{i+1} as described above.

Next, we show that using this guessing strategy, the pursuer efficiently locates the evader.

Lemma 1: At each round, if the pursuer follows the guessing strategy described above, he can locate the evader with probability at least $(1/N)$, where N is the total number of leaves of the triangulation tree T .

Proof: The lemma is proven by induction on the height of T . The basis, where the height is zero, corresponds to the case where the input polygon is a triangle. The pursuer trivially locates the evader with probability 1 in this case.

Let $p(\Delta)$ be the probability that the evader is located within a round, after the pursuer visits the triangle Δ . We inductively assume that the lemma is true for all trees of height less than or equal to i .

Given a triangulation tree of height $i + 1$, the probability of success starting from the root Δ_0 is

$$p(\Delta_0) \geq \min \left\{ \frac{l_0^1}{N} p(\Delta_0^1), \dots, \frac{l_0^k}{N} p(\Delta_0^k) \right\}. \quad (1)$$

²Note that the pursuer has unit speed.

³When the game starts, if the pursuer is located at a nonleaf triangle, he moves to an arbitrary leaf before starting the first round.

TABLE I
PURSUER'S STRATEGY FOR LOCATING THE EVADER. PLEASE REFER
TO FIG. 4 AND LEMMA 1 FOR NOTATION

LocateTheEvader(T : a triangulation of the environment)
Go to an arbitrary leaf triangle (<i>initialization</i>)
while the evader is not found
$\Delta_0 \leftarrow$ current triangle of the pursuer
$T \leftarrow T$ rooted at Δ_0
$i \leftarrow 0$
repeat
$C_i \leftarrow \{\Delta_i^j : \Delta_i^j \text{ is a child of } \Delta_i \text{ in } T\}$
$\Delta_{i+1} \leftarrow$ randomly chosen triangle from C_i where
Δ_i^j is chosen with probability $\frac{l_i^j}{\sum_j l_i^j}$
move from Δ_i to Δ_{i+1}
$i \leftarrow i + 1$
until Δ_i is a leaf triangle

Note that for all j , the subtrees rooted at the immediate children Δ_i^j have height at most i , therefore, by the inductive hypothesis, we have $p(\Delta_i^j) \geq (1/l_i^j)$ for all j , and the lemma follows. ■

Clearly, the number of leaves of any triangulation tree is less than the number of vertices of the polygon, therefore, at each round, the evader is located with probability at least $(1/n)$. Moreover, since the length of a round is $\text{diam}(P)$, we have the main result of this section in the following theorem.

Theorem 2: In any simply connected polygonal environment P , against any evader strategy, the expected time to locate the evader with a single pursuer is at most $n \cdot \text{diam}(P)$, where n is the number of vertices and $\text{diam}(P)$ is the diameter of the polygon.

The strategy for finding the evader is presented in Table I.

Remark 1: Any simply connected polygon can be partitioned into a minimum number of disjoint convex polygons in polynomial time [19], [20]. The dual of such a partition will also be a tree. Therefore, instead of using a triangulation dual, the pursuer can execute the strategy described above using the dual of the convex partition. However, in general, this does not improve the expected capture time. For example, for the polygon shown in Fig. 5, the number of leaves of the triangulation dual is equal to the number of leaves of the dual of a minimum convex partition.

Remark 2—Multiply Connected Environments: The strategy presented in this section requires the triangulation dual to be a tree, and therefore, it does not hold for multiply connected environments. To obtain an upper bound on the number of pursuers required for deterministic strategies, in [12] the following technique is presented. For an environment with h holes, $O(\sqrt{h})$ pursuers are used to reduce the environment into simply connected components. An additional $O(\log n)$ pursuers are used to deterministically clear each simply connected component. This yields an upper bound of $O(\sqrt{h} + \log n)$ pursuers. Using the same technique, we can establish an improved $O(\sqrt{h} + 1)$ bound for randomized strategies. Using $O(\sqrt{h})$ pursuers, we partition the environment into K simply connected polygons P_1, \dots, P_K . After partitioning the environment, we use an extra pursuer to locate the evader. The strategy of this pursuer is as follows. Pick a simply connected polygon \tilde{P} from among P_1, \dots, P_K , uniformly at random. Execute one round of the randomized strategy on \tilde{P} . The probability of success is easily seen




Fig. 5. For any randomized pursuer strategy, the expected time to capture the evader in this star with hooks is $O(n \cdot \text{diam}(P))$.

to be $(1/K) \cdot (1/n) \geq (1/n^2)$, and therefore, the expected time to capture the evader is $n^2 \cdot \text{diam}(P)$.

B. Lower Bounds

One might suspect that the expected time to locate an evader can be improved using a more sophisticated strategy. Unfortunately, this is not possible. The polygon in Fig. 5 is a k -star with hooks attached at the end of each spike (in the figure, $k = 8$). The evader's strategy is to choose a hook at random and hide there until the end of the game. In order to locate the evader, the expected number of spikes searched by the pursuer is $(k/2)$, and it takes $\text{diam}(P)$ steps to travel from one spike to another. Since the number of vertices is a constant multiple of k , the time it takes to locate the evader is $\Omega(n \cdot \text{diam}(P))$. In fact, using the well-known technique of Yao, this argument can be extended to show that the expected time to locate the evader for *any* randomized pursuer strategy is $\Omega(n \cdot \text{diam}(P))$ (see [15] for details).

We present the results of a simulation of the pursuer's and evader's strategy for such an environment in Fig. 6. For the simulation, the visibility-based pursuit-evasion game was played 1000 times. The average number of rounds for locating the evader is 8.9. This is in agreement with *Lemma 1*, since the number of the leaves of the triangulation dual is nine.

On the other hand, the randomized strategy may not be optimal for some environments. The simplest example of such an environment is a star-shaped polygon, such as the one shown in Fig. 7. In this environment, the optimal strategy is to go to a point (e.g., x in the figure) from where the entire polygon (hence, the evader) will be visible.

III. CAPTURING THE EVADER WITH TWO PURSUERS

In this section, we move on to the more challenging task of capturing the evader, defined as moving to the same point as the evader. We start by presenting a pursuit strategy for two pursuers (who can communicate at all times) to capture the evader. Later, we will show how to modify this strategy to obtain a strategy for two pursuers who have only line-of-sight (LOS) communication, and a single pursuer (at the expense of increasing the expected capture time).

The strategy of one of the pursuers is based on the solution to a problem known as the *lion and man* problem [2]. We present an

extension of this strategy in the case of a (possibly nonconvex) polygonal environment. One of the major difficulties for our pursuers is that the evader may not be visible at all times, in which case, the lion's strategy is not well defined. The second pursuer will use the strategy presented in the previous section to tackle this difficulty.

We start with a review of the lion's strategy.

A. Lion and Man Problem

The *lion and man* problem with discrete time in the nonnegative quadrant of the plane is attributed to Gale [21]. Let the initial positions of the lion and man be $L_0 = (x_0, y_0)$ and $M_0 = (x'_0, y'_0)$, respectively. In each round, first the man moves to any point in the quadrant at distance at most one from his current position, and then the lion does the same. The lion wins if he moves to the current position of the man. The man wins if he can keep escaping for infinitely many rounds. In [2], Sgall proves that when both $x'_0 < x_0$ and $y'_0 < y_0$, the lion always catches the man in a finite number of rounds (in remaining cases, the man wins the game). The number of moves required is bounded by a quadratic function in x_0, y_0 and the slope (or its inverse) of the line segment L_0M_0 .

B. Lion's Strategy

Let the initial positions of the lion and man be $L_0 = (x_0, y_0)$ and $M_0 = (x'_0, y'_0)$, respectively. In the beginning of the game, the lion finds a point C on the line M_0L_0 , such that L_0 is inside the segment M_0C and the circle with center C , radius $|CL_0|$, and passing through L_0 intersects both axes. Among all possible such circles, it chooses the one whose center is closest to the origin. C remains fixed throughout the game.

Let L and M denote the current positions of the lion and the man, respectively (see Fig. 8). Let M' denote the point the man moves to, $|MM'| \leq 1$. If $|LM'| \leq 1$, the lion catches the man. Otherwise, it moves to a point L' on the line $M'C$, such that $|L'L| = 1$. There are two such points, and it chooses the one closer to the man.

Definition 5: We will refer to this move as the *lion's move* from L with respect to C and M' (Fig. 8).

The lion's move maintains the following lemma.

Lemma 6 [2]: If the lion does not catch the man in the current move, then:

- 1) M' has both coordinates strictly smaller than C ;
- 2) L' is inside the segment $M'C$;
- 3) $|L'C|^2 \geq 1 + |LC|^2$.

Proof: See [2]. ■

C. Strategy to Capture the Evader

Let $p_1(t), p_2(t)$, and $e(t)$ denote the locations of the pursuers and the evader, respectively, at time t . In the beginning of the game, the two pursuers move together and search for the evader using the strategy described in the previous section. Without loss of generality, we assume that the game starts at $t = 0$, where $p_1(0) = p_2(0) = o$ and $e(0)$ is visible from o . We will sometimes refer to point o as the *origin*. The origin will be fixed until the evader is captured. Let $d_1(t) = d(p_1(t), o)$, $d_2(t) = d(p_2(t), o)$, and $d_e(t) = d(e(t), o)$.

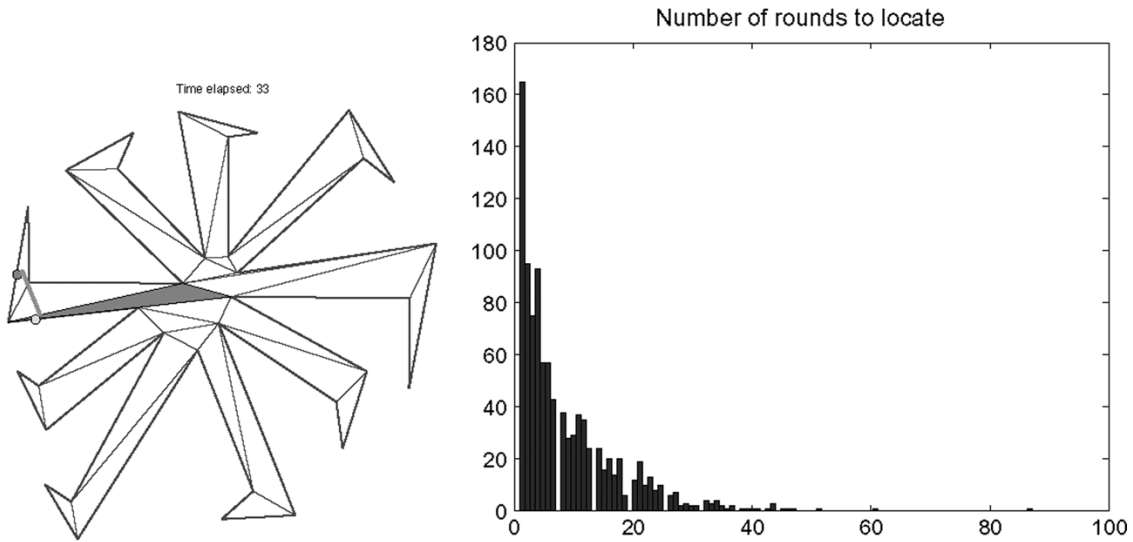


Fig. 6. **Left:** Instance of the simulator showing the triangulation of the environment, as well as the hiding location of the evader. **Right:** Histogram of the number of rounds required to locate evader in 1000 simulations. The mean μ and the standard deviation σ of the number of rounds was $\mu = 8.8960$ and $\sigma = 9.0479$.




Fig. 7. Star-shaped polygon. The optimal strategy is to go to a point (e.g., x in the figure) from where the entire polygon is visible.




Fig. 9. Pocket with respect to c and v .




Fig. 8. Lion's strategy.

Definition 7: Suppose $e(t)$ is visible from $p_1(t)$, but $e(t+1)$ is not visible from $p_1(t)$. This means that the shortest path P from $p_1(t)$ to $e(t+1)$ is composed of at least two line segments (*Property 1*). The first vertex on the path from $p_1(t)$ to $e(t+1)$ is called a *pseudo-blocking vertex*.

Let r be the ray starting from a vertex c and passing through another vertex v that is not adjacent to c . In the following, c will be the center of the circle for the lion's move, and v will be the pseudo-blocking vertex. Consider the first time the ray r leaves the polygon P after it passes through v , and let x be the point on $r \cap P$ just before this happens (see Fig. 9). The line segment vx splits the boundary of the polygon into two chains. The chain which does not contain the point c together with the line segment vx defines a polygon. We will refer to this polygon as *the pocket with respect to c and v* . The line segment vx is referred to as the *entrance of the pocket*.

We will use the following properties of pockets.

Property 2: Let α be a point on the line segment cv and β be a point in the pocket with respect to c and v . The line segment αv is contained in the shortest path from α to β (Fig. 9).

Property 3: Let R be a pocket with respect to c and v inside a polygon P . Any path from $\beta \in R$ to $\gamma \in P - R$ crosses the entrance of the pocket (Fig. 9).

Looking ahead, let us describe how we will use these properties. Suppose pursuer p_1 is moving toward the evader and the evader disappears. Let v be the current pseudo-blocking vertex. If p_1 moves toward v , *Property 2* implies that it is still moving




Fig. 10. Extended lion's move.

on the shortest path from the evader to the origin. If the evader becomes visible before p_1 reaches v , *Property 3* implies that it must cross the entrance of the pocket, and p_1 can continue its strategy (described in the next section) as if the evader has not disappeared.

If the evader is not visible when p_1 arrives at v , then v becomes a *blocking vertex*. At this point, the second pursuer will enter the game.

Next, we present the details of the strategies of p_1 and p_2 .

D. Strategy of Pursuer P_1

As stated earlier, we assume that pursuer p_1 is at least as fast as the evader. At time step t , p_1 moves according to the following strategy.

If the evader is visible, he performs an *extended lion's move*, which is defined as follows. Let τ be the shortest path from $e(t)$ to $e(t+1)$. Without loss of generality, p_1 will pretend that the evader followed τ . As a point x moves from $e(t)$ to $e(t+1)$ along τ , the vertices on the shortest path from x to the origin o may change. However, the number of changes is at most n . The first vertex on the shortest path from x to o must be one of the vertices of the polygon. Since τ is the shortest path from $e(t)$ to $e(t+1)$, each vertex of the polygon can be this first vertex for at most one contiguous subpath in τ . Let x_1, \dots, x_{k-1} correspond to the points on τ where such changes occur, we define $x_0 = e_t$ and $x_k = e(t+1)$. Let v_i be the first vertex on the shortest path from x_i to o . The extended lion's move consists of $k-1$ phases. During phase $i, i = 1, \dots, k$, pursuer p_1 performs the lion's move with respect to v_i and x_i (see Fig. 10). Note that the time spent by the pursuer in phase i is equal to the time spent by the evader in traveling from x_{i-1} from x_i .

If the evader was visible in the previous time step, but is not visible any more, let v be the pseudo-blocking vertex. Pursuer p_1 moves toward v until he reaches it. If the evader becomes visible before p_1 arrives at v , he continues with the lion's move. Otherwise, v becomes a blocking vertex.




Fig. 11. Summary of pursuer p_1 's strategy. **Left:** Evader is visible to p_1 , who proceeds with the extended lion's move. **Middle:** Evader disappears for the first time. It is known to be behind the pseudo-blocking vertex v . Pursuer p_1 moves toward v . **Right:** Pursuer p_1 cannot see the evader after arriving at v . The vertex v becomes a blocking vertex, and p_1 sends out p_2 to search for the evader. After finding the evader, p_2 will report v' , the new pseudo-blocking vertex.

If the evader is still not visible after p_1 reaches the blocking vertex, he waits for p_2 to report the location of the evader. Let R be the current pocket defined with respect to the blocking vertex and the current center. There are two possibilities.

- 1) The evader reveals itself to p_1 . Then, by *Property 3*, this must be while the evader is crossing the entrance of R . In this case, p_1 continues the game with the lion's move.
- 2) Pursuer p_2 finds the evader located at e . Let v' be the first vertex on the shortest path from v to e , and R' be the pocket with respect to v and v' . In this case, v' becomes a pseudo-blocking vertex, R' becomes the new pocket, and p_1 continues this strategy by moving toward v' .

An illustration of different modes of the pursuer's strategy is presented in Fig. 11.

E. Strategy of Pursuer p_2

The task of pursuer p_2 is to search for the evader when it is not visible to p_1 . When the evader disappears from the sight of p_1 , pursuer p_2 waits until p_1 reaches the blocking vertex. Afterwards, p_2 locates the evader using the strategy described in the previous section, and reports the location of the evader to p_1 .

F. Properties of Pursuer p_1 's Strategy

Lemma 8: For all times t , pursuer p_1 maintains the following invariants until the evader is caught:

- I1) $p_1(t)$ is on the shortest path from o to $e(t)$;
- I2) $d_1(t+1)^2 \geq d_1(t)^2 + (1/n)$ if $p_1(t) \neq p_1(t+1)$.

Proof of Invariant II: We prove the invariant by induction. Assume that it holds at time t .

First, consider the case where p_1 can see e at time t . Let the first vertex on the shortest path from $e(t)$ to o be u . It follows that $p(t)$ is in the line segment joining u to $e(t)$, since if $p(t)$ is between o and u on the shortest path, it would not be able to see $e(t)$.

Let x denote the evader's position at an arbitrary time in the time interval $[t, t+1)$. Suppose when the evader is at x , the first vertex on the shortest path from x to o changes from u to v . Note that p_1 can see the evader until this point. Then, the shortest path from x to o passing through u , and the shortest path from x

to o passing through v , have the same length. This implies that u, v , and x have to be collinear. For otherwise, a shorter path from x to o can be found in the interior of the polygon formed by these two presumed shortest paths from x to o , which is a contradiction.

This implies that either u is an ancestor or a descendant of v in the shortest-path tree rooted at o . If u is an ancestor, at the point x where the switch occurs, p_1 could either be on the segment vx , in which case it can continue the lion's move in the next phase, or p_1 is on the segment uv , in which case e will become invisible to p_1 after x . In this case, p_1 must be either moving toward a pseudo-blocking vertex or waiting at a blocking vertex. In both cases, the invariant is maintained by *Property 2*. If u is a descendant of v , then p_1 is already on the segment ux , and hence, on the segment vx . Hence, it can continue the lion's move in the next phase. The invariant is therefore maintained as a corollary of *Lemma 6*.

Otherwise, if p_1 does not see the evader at time t , he must be either waiting at a blocking vertex or moving toward a pseudo-blocking vertex. In both cases, the invariant is maintained by *Property 2*. ■

Proof of Invariant I2: If p_1 is moving toward a pseudo-blocking vertex, his distance to the origin is increasing by one, and the invariant is maintained.

Next, we show that the extended lion's move maintains the invariant. Suppose the lion's move has $k \leq n$ phases, and consider phase i of the extended lion's move, where the evader moves from the point x_{i-1} to x_i . Suppose that during this phase, pursuer p_1 moved from point y_{i-1} to y_i (see Fig. 10), and let v_i be the center of the circle for the lion's move during this phase.

Let $\omega_i = d(o, y_i) - d(o, y_{i-1})$.

As a corollary of *Lemma 6*, we have

$$d(y_i, o)^2 \geq d(y_{i-1}, o)^2 + \omega_i^2.$$

Summing up over all phases, we get the total progress as $\sum_{i=1}^k \omega_i^2$.

This expression, when subject to $\sum_{i=1}^k \omega_i = 1$, is minimized when all $\omega_1 = \dots = \omega_k = (1/k)$. Therefore, we have $d_1(t+1)^2 \geq d_1(t)^2 + (1/k)$, which implies the invariant I2. ■

The combined strategy of the two pursuers can be viewed as follows. Pursuer p_1 moves only when it knows the shortest path from the evader to the origin o . Performing the lion's move is equivalent to growing a disk inside the polygon, whose center is at the origin o and passes through the current location of p_1 . By invariant I1, the evader can never enter the disk. Further, the disk is still protected if p_1 does not move. Invariant I2 implies that, whenever p_1 moves, the disk monotonically grows, and the evader is eventually squeezed between p_1 and the polygon boundary.

Pursuer p_2 moves only when p_1 does not know the evader's path to the origin. It locates the evader using the randomized strategy given in the previous section, and reports its location to p_1 so that p_1 , in turn, can keep growing the disk and eventually capture the evader.

G. Expected Time to Capture

Let $T_1 = \text{diam}(P)$ be the time it takes pursuer 1 (who performs the lion's move) to travel the diameter of the polygon. By

Invariant I2 (*Lemma 8*), this pursuer will capture the evader in nT_1^2 steps. However, in the meantime, pursuer 2 may have to search for the evader.

The number of searches is bounded by the number of vertices. This is because once a vertex becomes a blocking vertex, it can never become a blocking vertex again. Next, we bound the length of each search. Recall that the probability of capturing the evader within a round is at least $(1/n)$ (*Lemma 1*). Using the inequality $(1+x) \leq e^x$, it can be easily shown that after $2n \ln n$ rounds, the probability of not finding the evader is at most $(1/n^2)$. Using the union bound, the probability of failure in any of the n searches is bounded by $n \cdot (1/n^2) = (1/n)$. Therefore, with probability $1 - (1/n)$, all n searches finish in total time $T_2 \cdot 2 \cdot n^2 \cdot \ln n$ with high probability, where T_2 is the time for pursuer p_2 to travel the diameter of the polygon.

In conclusion, the expected time to capture the evader is $O(nT_1^2 + T_2 \cdot (n^2 \ln n))$ with probability arbitrarily close to one.

Our main result is summarized by the following theorem.

Theorem 9: In any simply connected polygon, two pursuers p_1 and p_2 can capture an evader (whose speed is bounded by the speed of p_1) with probability arbitrarily close to one.

IV. EXTENSIONS OF THE TWO-PURSUER STRATEGY

In this section, we present three extensions of the two-pursuer strategy presented in the previous section. In Section IV-A, we show how a single pursuer can implement the same strategy at the expense of increased capture time. In Section IV-B, we show how the global communication requirement can be relaxed. Finally, in Section IV-C, we show that two pursuers can capture the evader, even if the polygon has a door through which the evader can escape and win the game.

A. Capturing the Evader With a Single Pursuer

Suppose we have only pursuer p_1 . In this case, instead of waiting for p_2 to find the evader, p_1 can guess the first vertex on the shortest path from the evader to its current location and move there.

Consider the shortest-path tree T from the origin o to the vertices of the polygon. For each vertex v , let $l(v)$ be the number of leaves of the subtree $T(v)$ of T rooted at the vertex v . Then the probability that the pursuer's guess will be successful if it is located at v is at least $(1/l(v))$. If the guess is correct and the evader is visible, the pursuer continues with the lion's move. However, in the case of a wrong guess, the evader may end up in an advantageous location and move toward the origin o , in which case, the pursuer must restart the game. Further, if all the guesses are correct, no vertex can be a blocking vertex more than once. Continuing this way, we can obtain a worst-case lower bound on the probability of success. Unfortunately, this bound can be possibly exponentially small in the number of reflex vertices in the environment. However, the expected time to capture the evader is still finite for any simply connected environment, and this strategy may still be practical for simple settings.

One might suspect that an analysis similar to the one in Section II can be applied to prove that the expected time to capture is polynomial. The reason such an analysis does not apply directly




Fig. 12. Pursuer p_1 is waiting at the current blocking vertex v . Pursuer p_2 finds the evader in the pocket with respect to v and u (the shaded region). The vertex u will become the new pseudo-blocking vertex.

is that even if the pursuer and the evader are co-located in a leaf triangle, the capture game still continues, and the evader can move to another triangle in the tree. Therefore, the number of guesses may exceed the depth of the tree, resulting in a possibly exponential capture time. This poses an interesting tradeoff between the pursuer's visibility and the capture time. If the pursuer can somehow track the evader at all times (perhaps using a satellite), then Lemma 8 implies that it could capture the evader in time $O(n \cdot \text{diam}(P)^2)$. If this is not possible, though, it can either use a second pursuer for locating the evader and still capture it in polynomial time, or simultaneously search and capture, which results in a much longer capture time.

B. Relaxing the Global Communication Requirement

When pursuer 1 arrives at a blocking vertex v , suppose the evader is not visible from v . Therefore, pursuer 2 starts searching for the evader, and finds it at time t .

In the previous section, we assumed that the two pursuers can communicate all the time. Hence, after finding the evader at time t , pursuer 2 can compute the first vertex u on the shortest path from $p_1(t)$ to $e(t)$, and report u to pursuer 1. Afterwards, u becomes the pseudo-blocking vertex and pursuer 1 starts moving toward u (Fig. 12).

The two pursuers can implement this strategy even if they have only LOS communication. That is, the pursuers can communicate if and only if they can see each other. The only modification required in the strategy is the following.

After finding the evader and computing the pseudo-blocking vertex u at time t , if pursuers 1 and 2 can not see each other, pursuer 2 can not report the vertex u (see Fig. 12). In this case, pursuer 2 starts moving toward pursuer 1 along the shortest path from $p_2(t)$ to $p_1(t)$. Let R be the pocket with respect to the current blocking vertex v and the pseudo-blocking vertex u (the shaded area in Fig. 12). Since pursuer 1 will eventually see pursuer 2, there are two possibilities, based on whether the evader becomes visible to pursuer 1 in the meantime.

If pursuer 1 sees the evader before seeing pursuer 2, this means that the evader is leaving R (Property 3). Further, at this time, pursuer 1 will be on the evader's shortest path to the origin, by Property 2. Therefore, in this case, pursuer 1 proceeds with

the extended lion's move, maintaining both invariants $I1$ and $I2$.

Otherwise, if pursuer 1 sees pursuer 2 first, the evader must be in the pocket with respect to v and u . In this case, pursuer 2 reports the new pseudo-blocking vertex u , and pursuer 1 proceeds as before by moving toward u . Note that pursuer 2 reports only the pseudo-blocking vertex, not the precise location of the evader.

C. Polygonal Rooms With a Door

In [13], Lee *et al.* studied the following variant of the pursuit-evasion problem. The input is a pair (P, d) where P is the polygonal room the game is played in, and d is a door, a point marked on the boundary of P . The goal is to devise a strategy for the pursuer to eventually see the evader, in such a way that the evader cannot escape through the door. The authors presented a characterization of polygons where a single pursuer with very narrow visibility (represented by a single ray) can locate the evader before it reaches the door.

In a similar scenario, the two-pursuer algorithm presented in Section III can be used to capture an evader before it exits through the door. The only modification necessary is the following. Initially, pursuer p_1 is located at the door d and waits until pursuer p_2 locates the evader. Afterwards, it continues with the lion's move with respect to d . This ensures that the evader can never enter the disk whose origin is d , and passes through the current location of p_1 . Therefore, the door is always protected until the evader is captured.

V. CONCLUSION AND FUTURE WORK

In this paper, we studied the visibility-based pursuit-evasion game, and showed that by using a randomized strategy, a single pursuer can locate an unpredictable evader in any simply connected polygonal environment. The evader may be arbitrarily faster than the pursuer, and it may know the location of the pursuer at all times.

The randomized strategy has some desirable properties. First, as shown in [12], there are polygonal environments which require an arbitrary number of pursuers if they are restricted to deterministic strategies. Therefore, in such environments, a randomized strategy is mandatory for locating the evader with a single pursuer. Moreover, even if the polygon is deterministically searchable by a single pursuer, it is known that some of these polygons require revisiting parts of the polygon $\Omega(n)$ times [12]. In such polygons, the expected time-to-capture with a randomized strategy is comparable to the time-to-capture with a deterministic strategy. However, the randomized strategies may be preferable to the deterministic strategies, as they do not require complicated data structures and costly preprocessing.

Second, the randomized strategy to locate the evader does not require an exact map of the environment. It is based on the dual graph of the triangulation, and as long as the structure of the triangulation dual is preserved, it is insensitive to errors in the map of the environment. An interesting research direction is to incorporate the navigation strategies in [7] which require a minimal representation of the environment.

Another interesting extension is the case of nonpolygonal environments. The randomized strategy can be used to locate the evader in nonpolygonal, simply connected environments. For example, this could be done by replacing the triangulation tree (Lemma 1) with the decomposition studied in [22]. Another alternative is to use the medial axis (cf. [16] for a definition and related properties) of a simply connected polygon instead of the triangulation tree.

We have also studied the more challenging problem of capturing the evader. For this problem, we presented a strategy for two pursuers (one of which is as fast as the evader) to capture the evader in an expected time polynomial in the number of vertices and the diameter of the environment. The strategy can be modified for a single pursuer, however, it is not clear whether the expected time-to-capture remains a polynomial in the number of vertices. We leave this as a future research direction.

REFERENCES

- [1] V. Isler, S. Kannan, and S. Khanna, "Locating and capturing an evader in a polygonal environment," in *Proc. 6th Workshop Algorithmic Problems Robot.*, 2004, pp. 351–367.
- [2] J. Sgall, "Solution of David Gale's lion and man problem," *Theoret. Comput. Sci.*, vol. 259, no. 1-2, pp. 663–670, 2001.
- [3] J. Hespanha, H. J. Kim, and S. Sastry, "Multiple-agent probabilistic pursuit-evasion games," in *Proc. 38th Conf. Decision Contr.*, 1999, pp. 2432–2437.
- [4] R. Vidal, O. Shakernia, J. Kim, D. Shim, and S. Sastry, "Probabilistic pursuit-evasion games: Theory, implementation, and experimental evaluation," *IEEE Trans. Robot. Autom.*, vol. 18, no. 5, pp. 662–669, Oct. 2002.
- [5] J. P. Hespanha, G. J. Pappas, and M. Prandini, "Greedy control for hybrid pursuit-evasion games," in *Proc. Eur. Control Conf.*, Porto, Portugal, Sep. 2000, pp. 2621–2626.
- [6] A. Sarmiento, R. Murrieta-Cid, and S. Hutchinson, "An efficient strategy for rapidly finding an object in a polygonal world," in *Proc. IEEE/RSJ Int. Conf. Intell. Robots Optim. Syst.*, 2003, pp. 1153–1158.
- [7] B. Tovar, S. M. LaValle, and R. Murrieta, "Optimal navigation and object finding without geometric maps or localization," in *Proc. IEEE Int. Conf. Robot. Autom.*, 2003, pp. 464–470.
- [8] R. Isaacs, *Differential Games*. New York: Dover, 1965.
- [9] T. Basar and G. J. Olsder, *Dynamic Noncooperative Game Theory*. Philadelphia, PA: SIAM, 1998.
- [10] I. Suzuki and M. Yamashita, "Searching for a mobile intruder in a polygonal region," in *SIAM J. Comput.*, vol. 21, 1992, pp. 863–888.
- [11] S.-M. Park, J.-H. Lee, and K.-Y. Chwa, "Visibility-based pursuit evasion in a polygonal region by a searcher," presented at *Int. Colloq. Automata, Lang. Program.*, [Online]. Available: citeseer.nj.nec.com/park01visibilitybased.html
- [12] L. J. Guibas, J.-C. Latombe, S. M. LaValle, D. Lin, and R. Motwani, "A visibility-based pursuit-evasion problem," *Int. J. Computat. Geom. Applicat.*, vol. 9, no. 4/5, pp. 471–1999, 1999.
- [13] J.-H. Lee, S.-M. Park, and K.-Y. Chwa, "Searching a polygonal room with one door by a 1-searcher," *Int. J. Computat. Geom. Applicat.*, vol. 10, no. 2, pp. 201–220, 2000.
- [14] S. M. LaValle, D. Lin, L. J. Guibas, J.-C. Latombe, and R. Motwani, "Finding an unpredictable target in a workspace with obstacles," in *Proc. IEEE Int. Conf. Robot. Autom.*, 1997, pp. 737–742.
- [15] R. Motwani and P. Raghavan, *Randomized Algorithms*. Cambridge, U.K.: Cambridge Univ. Press, 1995.
- [16] J. O'Rourke, *Computational Geometry in C*. Cambridge, U.K.: Cambridge Univ. Press, 1998.
- [17] J. E. Littlewood, *A Mathematician's Miscellany*. London, U.K.: Methuen, 1953.
- [18] L. Alonso, A. S. Goldstein, and E. M. Reingold, "'Lion and man': Upper and lower bounds," *ORSA J. Comput.*, vol. 4, no. 4, pp. 447–452, 1992.
- [19] M. Keil, "Decomposing a polygon into simpler components," *SIAM J. Comput.*, vol. 14, pp. 799–817, 1985.
- [20] B. Chazelle and D. Dobkin, "Optimal convex decompositions," in *Computational Geometry*, G. Toussaint, Ed. Amsterdam, The Netherlands: North-Holland, 1985, pp. 63–133.
- [21] R. Guy, "Unsolved problems in combinatorial games," in *Proc. Symp. Appl. Math.*, 1991, pp. 183–189.
- [22] S. M. LaValle and J. Hinrichsen, "Visibility-based pursuit evasion: The case of curved environments," *IEEE Trans. Robot. Autom.*, vol. 17, no. 2, pp. 196–201, Apr. 2001.



Volkan Isler (S'01) received the M.S.E. and Ph.D. degrees in computer and information science from the University of Pennsylvania, Philadelphia, and the B.S. degree in computer engineering from Bogazici University, Istanbul, Turkey.

He is currently a Postdoctoral Researcher with the Center for Information Technology Research in the Interest of Society (CITRIS), University of California, Berkeley. His research interests are in robotics (pursuit evasion, exploration, motion planning), sensor networks (deployment, target tracking,

and localization) and computer vision (tele-immersion, model reconstruction, and segmentation).



Sampath Kannan received the M.S. degree from Princeton University, Princeton, NJ, in 1984, and the Ph.D. degree from the University of California, Berkeley, in 1989, both in computer science.

He is currently a Professor with the Computer and Information Science Department, University of Pennsylvania, Philadelphia. His research interests are algorithms, computational biology, and formal methods for software reliability. He is currently an Editor of the *ACM Transactions on Algorithms*.



Sanjeev Khanna received the Ph.D. degree in computer science from Stanford University, Stanford, CA, and the B.S. degree in computer science from the Birla Institute of Technology, Pilani, India.

He is currently a Professor with the Computer and Information Science Department, University of Pennsylvania, Philadelphia. His research interests are in approximation algorithms, combinatorial optimization, computational biology, and databases.